

✓ Gesture Recognition

In this group project, you are going to build a 3D Conv model that will be able to predict the 5 gestures correctly. Please import the following libraries to get started.

```
1 import numpy as np
2 import os
3 import datetime
4 import cv2
5 import matplotlib.pyplot as plt
6 import warnings
7 warnings.filterwarnings('ignore')
```

We set the random seed so that the results don't vary drastically.

```
1 np.random.seed(30)
2 import random as rn
3 rn.seed(30)
4 from keras import backend as K
5 import tensorflow as tf
6 tf.random.set_seed(30)
```

In this block, you read the folder names for training and validation. You also set the batch_size here. Note that you set the batch size in such a way that you are able to use the GPU in full capacity. You keep increasing the batch size until the machine throws an error.

```
1 train_doc = np.random.permutation(open('/notebooks/storage/Final_data/Collated_training/train.csv').readlines())
2 val_doc = np.random.permutation(open('/notebooks/storage/Final_data/Collated_training/val.csv').readlines())
3 batch_size = 15 #experiment with the batch size
```

✓ Checking Data Imbalance

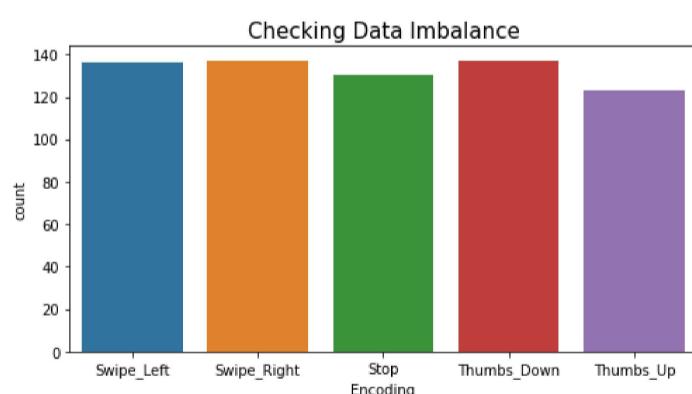
```
1 import pandas as pd
2 df = pd.read_csv('/content/train.csv', header=None)
3 df.columns = ['Action']
4 df['Encoding'] = df['Action'].apply(lambda x: x[-1])
5 data_count = df['Encoding'].value_counts()
6 data_count
```

```
1 137
3 137
0 136
2 130
4 123
Name: Encoding, dtype: int64
```

```
1 {x.split(';')[1]:x.split(';')[2][0] for x in train_doc}
```

```
1 {'Left_Swipe_new_Left_Swipe_new': '0',
 'Left_Swipe_new': '0',
 'Right_Swipe_new': '1',
 'Right_Swipe_new': '1',
 'Stop_Gesture_new': '2',
 'Stop_new': '2',
 'Thumbs_Down_new': '3',
 'Thumbs_Up_new': '4',
 'Thumbs_Down_new': '3',
 'Thumbs_Up_new': '4'}
```

```
1 import seaborn as sns
2 plt.figure(figsize=(8,4))
3 sns.countplot(df['Encoding'])
4 plt.title('Checking Data Imbalance', size=15)
5 plt.xticks(ticks=[0,1,2,3,4], labels=['Swipe_Left', 'Swipe_Right', 'Stop', 'Thumbs_Down', 'Thumbs_Up'])
6 plt.show()
```



✓ Generator

This is one of the most important part of the code. The overall structure of the generator has been given. In the generator, you are going to preprocess the images as you have images of 2 different dimensions as well as create a batch of video frames. You have to experiment with img_idx, y, z and normalization such that you get high accuracy.

✓ Using OpenCv instead of imread for the generator function due to following reasons.

- OpenCv is better in performance/speed of image processing as compared to skimage.
- OpenCv is more convenient to use than skimage.

```

1 def get_batch(source_path, folder_list, batch_size,batch,t,img_idx):
2     batch_data = np.zeros((batch_size,len(img_idx),120,120,3)) # x is the number of images you use for each video, (y,z) is the final size of the input images and 3 is the number of channels RGB
3     batch_labels = np.zeros((batch_size,5)) # batch_labels is the one hot representation of the output
4     for folder in range(batch_size): # iterate over the batch_size
5         imgs = os.listdir(source_path+'/'+ t[folder + (batch*batch_size)].split(';')[0]) # read all the images in the folder
6         for idx,item in enumerate(img_idx): # Iterate over the frames/images of a folder to read them in
7             image = cv2.imread(source_path+'/'+ t[folder + (batch*batch_size)].strip().split(';')[0]+"/"+imgs[item]).astype(np.float32)
8             h,w = image.shape[0:2]
9             #crop the images and resize them. Note that the images are of 2 different shape
10            #and the conv3D will throw error if the inputs in a batch have different shapes
11            if (h>120) and (w>160):
12                image = cv2.resize(image,(160,120),interpolation = cv2.INTER_NEAREST)
13
14            #Crop the Image
15            image = image[:,20:140,:]
16
17            batch_data[folder,idx,:,:,:0] = cv2.normalize(image[:, :, 0],None,0.00,1.00,cv2.NORM_MINMAX)#normalise and feed in the image
18            batch_data[folder,idx,:,:,:1] = cv2.normalize(image[:, :, 1],None,0.00,1.00,cv2.NORM_MINMAX)#normalise and feed in the image
19            batch_data[folder,idx,:,:,:2] = cv2.normalize(image[:, :, 2],None,0.00,1.00,cv2.NORM_MINMAX)#normalise and feed in the image
20
21            batch_labels[folder, int(t[folder + (batch*batch_size)].strip().split(';')[2])] = 1
22    return batch_data, batch_labels

```

```

1 def generator(source_path, folder_list, batch_size):
2     print( 'Source path = ', source_path, '; batch size =', batch_size)
3     img_idx = range(0,30)
4     while True:
5         t = np.random.permutation(folder_list)
6         num_batches = len(folder_list) // batch_size # calculate the number of batches
7         for batch in range(num_batches): # we iterate over the number of batches
8             yield get_batch(source_path, folder_list, batch_size,batch,t,img_idx)
9
10        # write the code for the remaining data points which are left after full batches
11        remaining = len(folder_list) % batch_size
12        if remaining != 0:
13            batch_size = remaining
14            yield get_batch(source_path, folder_list, batch_size,batch,t,img_idx)
15
16
17

```

```

1 # def generator(source_path, folder_list, batch_size):
2 #     print( 'Source path = ', source_path, '; batch size =', batch_size)
3 #     img_idx = #create a list of image numbers you want to use for a particular video
4 #     while True:
5 #         t = np.random.permutation(folder_list)
6 #         num_batches = # calculate the number of batches
7 #         for batch in range(num_batches): # we iterate over the number of batches
8 #             batch_data = np.zeros((batch_size,x,y,z,3)) # x is the number of images you use for each video, (y,z) is the final size of the input images and 3 is the number of channels RGB
9 #             batch_labels = np.zeros((batch_size,5)) # batch_labels is the one hot representation of the output
10 #             for folder in range(batch_size): # iterate over the batch_size
11 #                 imgs = os.listdir(source_path+'/'+ t[folder + (batch*batch_size)].split(';')[0]) # read all the images in the folder
12 #                 for idx,item in enumerate(img_idx): # Iterate over the frames/images of a folder to read them in
13 #                     image = imread(source_path+'/'+ t[folder + (batch*batch_size)].strip().split(';')[0]+"/"+imgs[item]).astype(np.float32)
14
15 #                     #crop the images and resize them. Note that the images are of 2 different shape
16 #                     #and the conv3D will throw error if the inputs in a batch have different shapes
17
18 #                     batch_data[folder,idx,:,:,:0] = #normalise and feed in the image
19 #                     batch_data[folder,idx,:,:,:1] = #normalise and feed in the image
20 #                     batch_data[folder,idx,:,:,:2] = #normalise and feed in the image
21
22 #                     batch_labels[folder, int(t[folder + (batch*batch_size)].strip().split(';')[2])] = 1
23 #     yield batch_data, batch_labels #you yield the batch_data and the batch_labels, remember what does yield do
24
25
26 #     # write the code for the remaining data points which are left after full batches
27

```

Note here that a video is represented above in the generator as (number of images, height, width, number of channels). Take this into consideration while creating the model architecture.

```

1 curr_dt_time = datetime.datetime.now()
2 # train_path = '/notebooks/storage/Final_data/Collated_training/train'
3 # val_path = '/notebooks/storage/Final_data/Collated_training/val'
4
5 # train_path = '/content/Project_data/train'
6 # val_path = '/content/Project_data/val'
7
8 train_path = '/content/train'
9 val_path = '/content/val'
10 num_train_sequences = len(train_doc)
11 print('# training sequences =', num_train_sequences)
12 num_val_sequences = len(val_doc)
13 print('# validation sequences =', num_val_sequences)
14 num_epochs = 20 # choose the number of epochs
15 print ('# epochs =', num_epochs)

→ # training sequences = 663
# validation sequences = 100
# epochs = 20

```

Function for the Batch Size

```

1 def gen_train_val (batch_size):
2     train_generator = generator(train_path, train_doc, batch_size)
3     val_generator = generator(val_path, val_doc, batch_size)
4
5     if (num_train_sequences%batch_size) == 0:
6         steps_per_epoch = int(num_train_sequences/batch_size)
7     else:
8         steps_per_epoch = (num_train_sequences//batch_size) + 1
9
10    if (num_val_sequences%batch_size) == 0:
11        validation_steps = int(num_val_sequences/batch_size)
12    else:
13        validation_steps = (num_val_sequences//batch_size) + 1
14
15    return train_generator,val_generator,steps_per_epoch,validation_steps

```

Function for Checkpoint

```
1 def set_checkpoint(model_name,fact=0.1,pati=8):
2
3     logdir = os.path.join("logs",model_name)
4     tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1,write_images=True)
5
6     filepath = os.path.join("model",model_name+'.h5')
7     checkpoint = ModelCheckpoint(filepath, monitor='val_loss', verbose=1, save_best_only=True, mode='auto')
8
9     LR = ReduceLROnPlateau(monitor='val_loss', factor=fact,patience=pati,mode='min')
10
11    return [tensorboard_callback, checkpoint, LR]
```

Function for train and validation generator

```
1 def gen_train_val (batch_size):
2     train_generator = generator(train_path, train_doc, batch_size)
3     val_generator = generator(val_path, val_doc, batch_size)
4
5     if (num_train_sequences%batch_size) == 0:
6         steps_per_epoch = int(num_train_sequences/batch_size)
7     else:
8         steps_per_epoch = (num_train_sequences//batch_size) + 1
9
10    if (num_val_sequences%batch_size) == 0:
11        validation_steps = int(num_val_sequences/batch_size)
12    else:
13        validation_steps = (num_val_sequences//batch_size) + 1
14
15    return train_generator,val_generator,steps_per_epoch,validation_steps
```

Model

Here you make the model using different functionalities that Keras provides. Remember to use Conv3D and MaxPooling3D and not Conv2D and Maxpooling2D for a 3D convolution model. You would want to use TimeDistributed while building a Conv2D + RNN model. Also remember that the last layer is the softmax. Design the network in such a way that the model is able to give good accuracy on the least number of parameters so that it can fit in the memory of the webcam.

```
1 from keras.models import Sequential, Model
2 from keras.layers import Dense, GRU, Flatten, TimeDistributed, Flatten, BatchNormalization, Activation
3 from keras.layers.convolutional import Conv3D, MaxPooling3D
4 from keras.layers import Conv2D,MaxPooling2D,Dropout
5 from keras.callbacks import ModelCheckpoint, ReduceLROnPlateau, TensorBoard
6 from keras import optimizers
7 %load_ext tensorboard
```

Convolution 3D : Initial Model

- Creating a base model where we just want to check if data fits the model.
- According to our expectation the model should overfit.

```
1 #write your model here
2 model = Sequential([
3     Conv3D(64, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform', input_shape=(30,120,120,3)),
4     MaxPooling3D(pool_size=(2, 2, 2)),
5     Conv3D(128, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
6     MaxPooling3D(pool_size=(2, 2, 2)),
7     Flatten(),
8     Dense(5, activation='softmax')
9 ])
```

Now that you have written the model, the next step is to compile the model. When you print the summary of the model, you'll see the total number of parameters you have to train.

```
1 #optimiser = tf.keras.optimizers.Adam(learning_rate=0.001) #write your optimizer
2 optimiser = tf.keras.optimizers.SGD(learning_rate=0.001)
3 model.compile(optimizer=optimiser, loss='categorical_crossentropy', metrics=['categorical_accuracy'])
4 print (model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv3d (Conv3D)	(None, 28, 118, 118, 64)	5248
<hr/>		
max_pooling3d (MaxPooling3D)	(None, 14, 59, 59, 64)	0
<hr/>		
conv3d_1 (Conv3D)	(None, 12, 57, 57, 128)	221312
max_pooling3d_1 (MaxPooling 3D)	(None, 6, 28, 28, 128)	0
<hr/>		
flatten (Flatten)	(None, 602112)	0
<hr/>		
dense (Dense)	(None, 5)	3010565
<hr/>		
Total params: 3,237,125		
Trainable params: 3,237,125		
Non-trainable params: 0		
<hr/>		
None		

Let us create the train_generator and the val_generator which will be used in .fit_generator.

```
1 train_generator,val_generator,steps_per_epoch,validation_steps = gen_train_val(batch_size=15)
```

The steps_per_epoch and validation_steps are used by fit_generator to decide the number of next() calls it need to make.

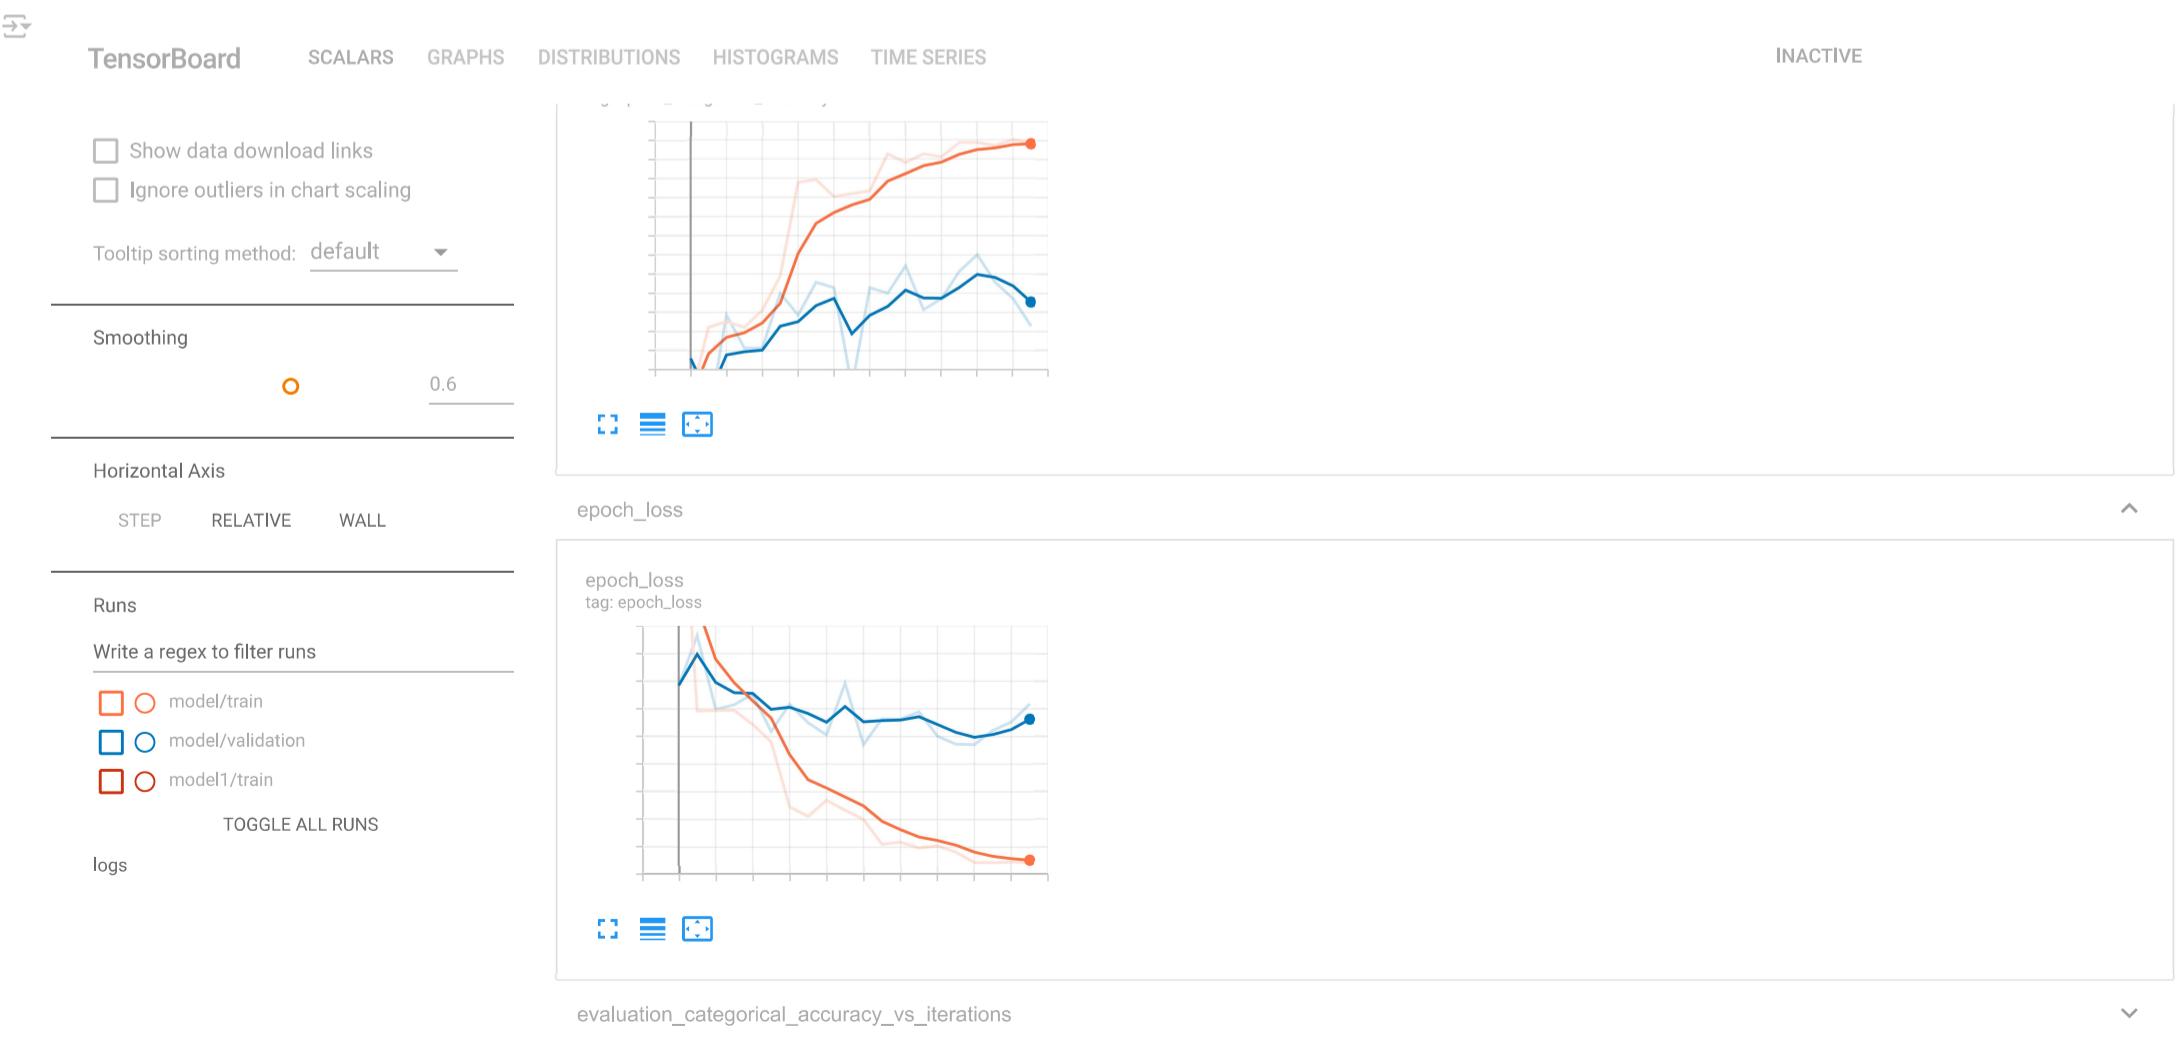
Let us now fit the model. This will start training the model and with the help of the checkpoints, you'll be able to save the model at the end of each epoch.

```
1 # Check Point for the base
2 callbacks_list = set_checkpoint('model')

1 model.fit(train_generator, steps_per_epoch=steps_per_epoch, epochs=num_epochs,
2             callbacks=callbacks_list, validation_data=val_generator,
3             validation_steps=validation_steps)

Epoch 00005: val_loss did not improve from 1.19525
45/45 [=====] - 12s 275ms/step - loss: 1.0876 - categorical_accuracy: 0.5556 - val_loss: 1.3051 - val_categorical_accuracy: 0.4571 - lr: 0.0010
Epoch 6/20
45/45 [=====] - ETA: 0s - loss: 0.9613 - categorical_accuracy: 0.6444
Epoch 00006: val_loss improved from 1.19525 to 1.03635, saving model to model/model.h5
45/45 [=====] - 13s 296ms/step - loss: 0.9613 - categorical_accuracy: 0.6444 - val_loss: 1.0363 - val_categorical_accuracy: 0.6000 - lr: 0.0010
Epoch 7/20
45/45 [=====] - ETA: 0s - loss: 0.4914 - categorical_accuracy: 0.8889
Epoch 00007: val_loss did not improve from 1.03635
45/45 [=====] - 13s 286ms/step - loss: 0.4914 - categorical_accuracy: 0.8889 - val_loss: 1.2330 - val_categorical_accuracy: 0.5429 - lr: 0.0010
Epoch 8/20
45/45 [=====] - ETA: 0s - loss: 0.4227 - categorical_accuracy: 0.8963
Epoch 00008: val_loss did not improve from 1.03635
45/45 [=====] - 13s 303ms/step - loss: 0.4227 - categorical_accuracy: 0.8963 - val_loss: 1.0997 - val_categorical_accuracy: 0.6286 - lr: 0.0010
Epoch 9/20
45/45 [=====] - ETA: 0s - loss: 0.5393 - categorical_accuracy: 0.8519
Epoch 00009: val_loss improved from 1.03635 to 1.01048, saving model to model/model.h5
45/45 [=====] - 13s 295ms/step - loss: 0.5393 - categorical_accuracy: 0.8519 - val_loss: 1.0105 - val_categorical_accuracy: 0.6143 - lr: 0.0010
Epoch 10/20
45/45 [=====] - ETA: 0s - loss: 0.4675 - categorical_accuracy: 0.8593
Epoch 00010: val_loss did not improve from 1.01048
45/45 [=====] - 13s 283ms/step - loss: 0.4675 - categorical_accuracy: 0.8593 - val_loss: 1.3852 - val_categorical_accuracy: 0.3571 - lr: 0.0010
Epoch 11/20
45/45 [=====] - ETA: 0s - loss: 0.4012 - categorical_accuracy: 0.8667
Epoch 00011: val_loss improved from 1.01048 to 0.94111, saving model to model/model.h5
45/45 [=====] - 12s 268ms/step - loss: 0.4012 - categorical_accuracy: 0.8667 - val_loss: 0.9411 - val_categorical_accuracy: 0.6143 - lr: 0.0010
Epoch 12/20
45/45 [=====] - ETA: 0s - loss: 0.2215 - categorical_accuracy: 0.9630
Epoch 00012: val_loss did not improve from 0.94111
45/45 [=====] - 13s 296ms/step - loss: 0.2215 - categorical_accuracy: 0.9630 - val_loss: 1.1275 - val_categorical_accuracy: 0.6000 - lr: 0.0010
Epoch 13/20
45/45 [=====] - ETA: 0s - loss: 0.2367 - categorical_accuracy: 0.9407
Epoch 00013: val_loss did not improve from 0.94111
45/45 [=====] - 12s 280ms/step - loss: 0.2367 - categorical_accuracy: 0.9407 - val_loss: 1.1249 - val_categorical_accuracy: 0.6714 - lr: 0.0010
Epoch 14/20
45/45 [=====] - ETA: 0s - loss: 0.1941 - categorical_accuracy: 0.9630
Epoch 00014: val_loss did not improve from 0.94111
45/45 [=====] - 13s 302ms/step - loss: 0.1941 - categorical_accuracy: 0.9630 - val_loss: 1.1787 - val_categorical_accuracy: 0.5571 - lr: 0.0010
Epoch 15/20
45/45 [=====] - ETA: 0s - loss: 0.2096 - categorical_accuracy: 0.9556
Epoch 00015: val_loss did not improve from 0.94111
45/45 [=====] - 12s 276ms/step - loss: 0.2096 - categorical_accuracy: 0.9556 - val_loss: 1.0027 - val_categorical_accuracy: 0.5857 - lr: 0.0010
Epoch 16/20
45/45 [=====] - ETA: 0s - loss: 0.1641 - categorical_accuracy: 0.9926
Epoch 00016: val_loss did not improve from 0.94111
45/45 [=====] - 12s 272ms/step - loss: 0.1641 - categorical_accuracy: 0.9926 - val_loss: 0.9444 - val_categorical_accuracy: 0.6571 - lr: 0.0010
Epoch 17/20
45/45 [=====] - ETA: 0s - loss: 0.0901 - categorical_accuracy: 0.9926
Epoch 00017: val_loss improved from 0.94111 to 0.94098, saving model to model/model.h5
45/45 [=====] - 13s 286ms/step - loss: 0.0901 - categorical_accuracy: 0.9926 - val_loss: 0.9410 - val_categorical_accuracy: 0.7000 - lr: 0.0010
Epoch 18/20
45/45 [=====] - ETA: 0s - loss: 0.0877 - categorical_accuracy: 0.9852
Epoch 00018: val_loss did not improve from 0.94098
45/45 [=====] - 13s 287ms/step - loss: 0.0877 - categorical_accuracy: 0.9852 - val_loss: 1.0434 - val_categorical_accuracy: 0.6286 - lr: 0.0010
Epoch 19/20
45/45 [=====] - ETA: 0s - loss: 0.0926 - categorical_accuracy: 1.0000
Epoch 00019: val_loss did not improve from 0.94098
45/45 [=====] - 13s 293ms/step - loss: 0.0926 - categorical_accuracy: 1.0000 - val_loss: 1.1048 - val_categorical_accuracy: 0.5857 - lr: 0.0010
Epoch 20/20
```

```
1 # Visualizing the Model Results using TensorBoard
2
3 %tensorboard --logdir logs
```



Best Epoch (Epoch 17) Result from the above odel Train Accuracy : 99.26 Validation Accuracy : 70.00

- We can clearly see that the above Model is overfitting, it was the purpose of the initial model to verify the fit

✓ Convolution 3D Models

- Starting with the CONV 3D models
- We see the accuracy by playing around with the Batch size, Optimizer, model architecture
- We will also experiment with the no. of layers and Drop outs in the Model Architecture

✓ Conv 3D Model1 : (Batch Size: 15 , No. of Frames = 30 , Optimizer=SGD)

```

1 #Changing the batch Size
2 train_generator,val_generator,steps_per_epoch,validation_steps = gen_train_val(batch_size=15)
3
4 #Setting the Callback Parameters
5 callbacks_list1 = set_checkpoint('model1',pati=6)

1 model1 = Sequential([
2     Conv3D(16, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform', input_shape=(30,120,120,3)),
3     BatchNormalization(),
4     Conv3D(32, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
5     BatchNormalization(),
6     Conv3D(64, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
7     BatchNormalization(),
8     MaxPooling3D(pool_size=(2, 2, 2)),
9     Conv3D(128, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
10    BatchNormalization(),
11    MaxPooling3D(pool_size=(2, 2, 2)),
12    Flatten(),
13    Dense(64, activation='relu', kernel_initializer='he_uniform'),
14    BatchNormalization(),
15    Dense(5, activation='softmax')
16 ])

```

```

1 optimiser = tf.keras.optimizers.SGD(learning_rate=0.001) # write your optimizer
2 model1.compile(optimizer=optimiser, loss='categorical_crossentropy', metrics=['categorical_accuracy'])
3 print (model1.summary())

```

→ Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv3d_2 (Conv3D)	(None, 28, 118, 118, 16)	1312
batch_normalization (BatchN ormalization)		
conv3d_3 (Conv3D)	(None, 26, 116, 116, 32)	13856
batch_normalization_1 (BatchN hNormalization)	(None, 26, 116, 116, 32)	128
conv3d_4 (Conv3D)	(None, 24, 114, 114, 64)	55360
batch_normalization_2 (BatchN hNormalization)	(None, 24, 114, 114, 64)	256
max_pooling3d_2 (MaxPooling 3D)	(None, 12, 57, 57, 64)	0
conv3d_5 (Conv3D)	(None, 10, 55, 55, 128)	221312
batch_normalization_3 (BatchN hNormalization)	(None, 10, 55, 55, 128)	512
max_pooling3d_3 (MaxPooling 3D)	(None, 5, 27, 27, 128)	0
flatten_1 (Flatten)	(None, 466560)	0
dense_1 (Dense)	(None, 64)	29859904
batch_normalization_4 (BatchN hNormalization)	(None, 64)	256
dense_2 (Dense)	(None, 5)	325
=====		
Total params:	30,153,285	
Trainable params:	30,152,677	
Non-trainable params:	608	

None

```

1 model1.fit(train_generator,
2     steps_per_epoch=steps_per_epoch,
3     epochs=50,
4     verbose=1,
5     validation_data=val_generator,
6     validation_steps=validation_steps,
7     callbacks=callbacks_list1)

```

→

```

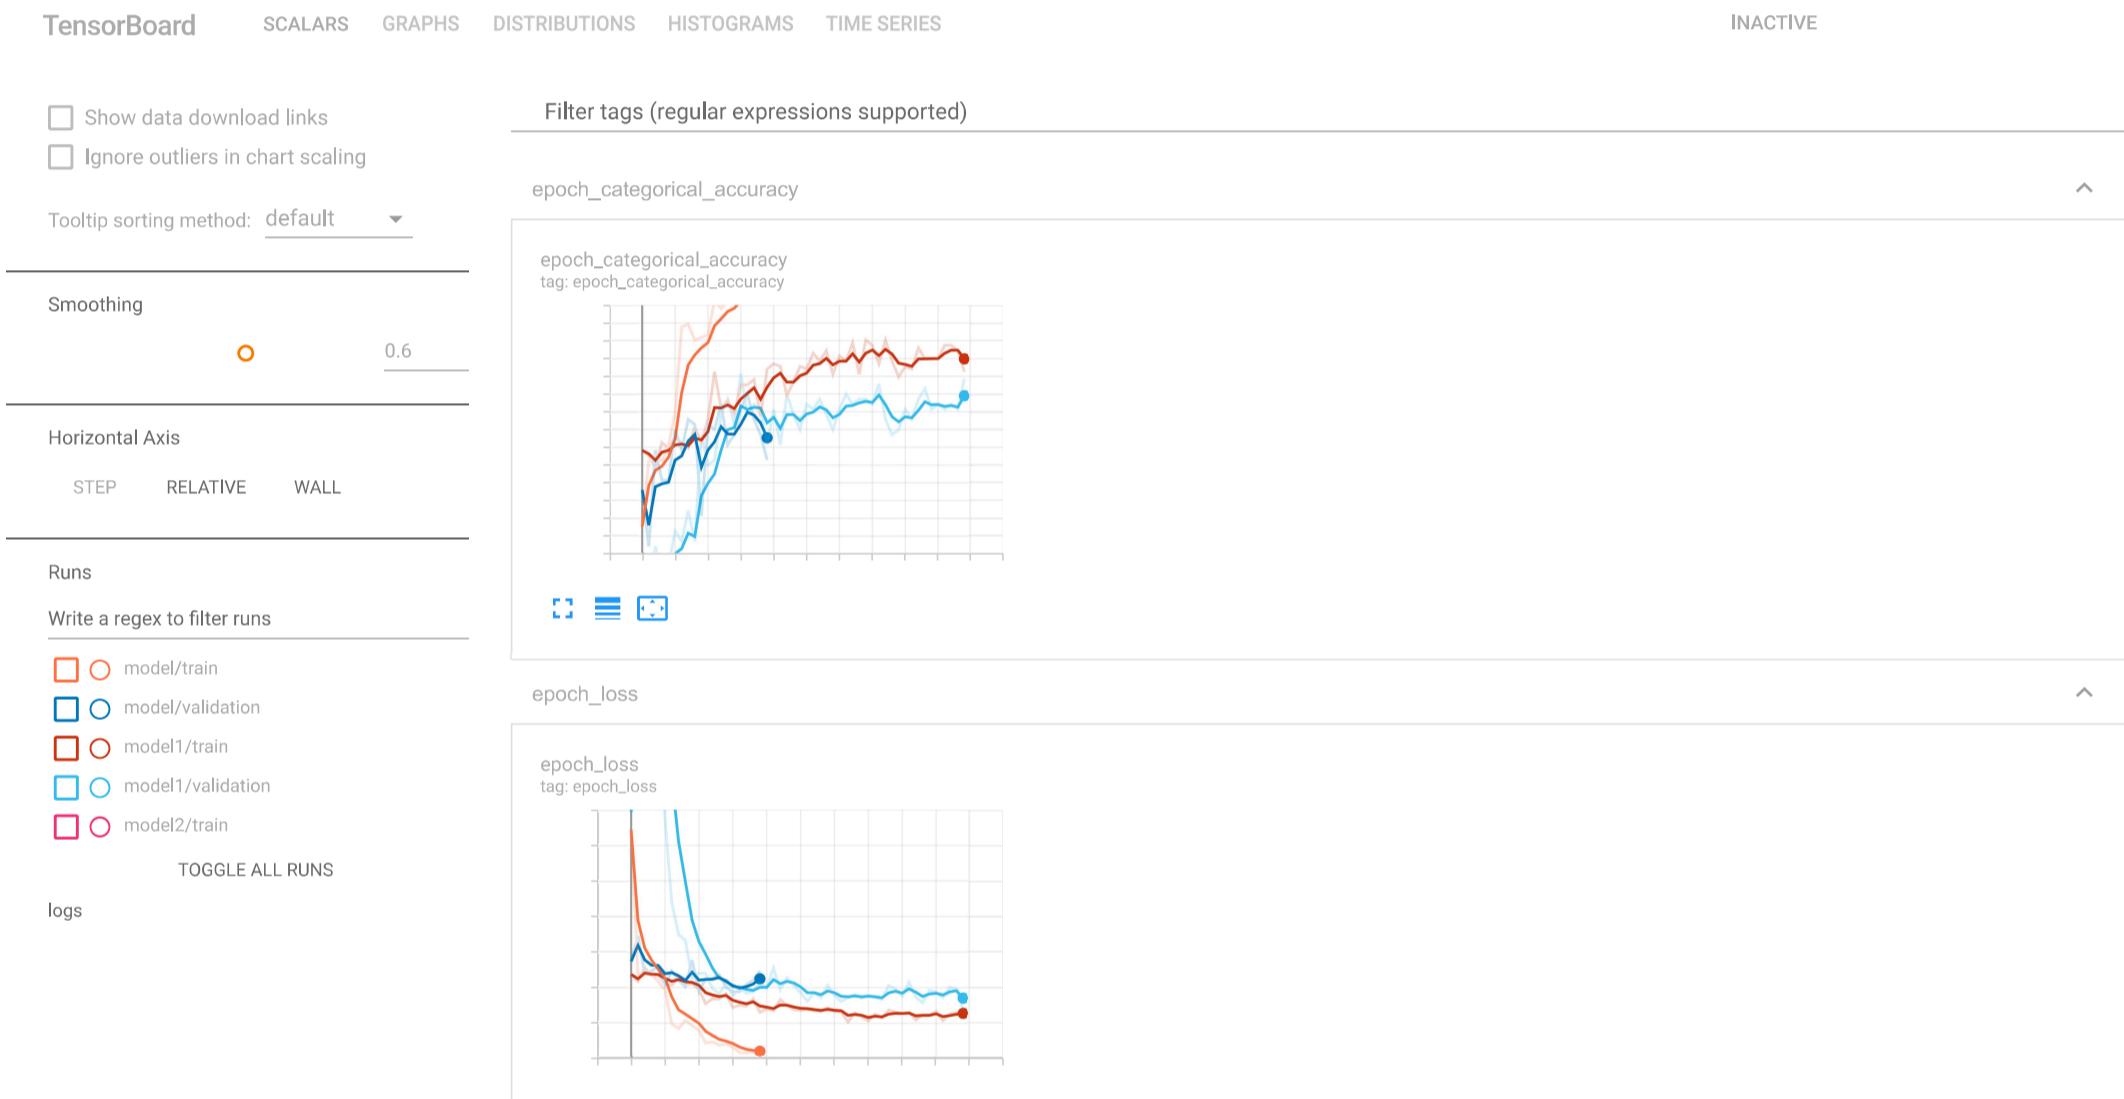
Epoch 00035: val_loss did not improve from 0.80483
45/45 [=====] - 22s 487ms/step - loss: 0.5936 - categorical_accuracy: 0.8519 - val_loss: 0.8523 - val_categorical_accuracy: 0.6857 - lr: 1.0000e-04
Epoch 36/50
45/45 [=====] - ETA: 0s - loss: 0.5338 - categorical_accuracy: 0.8370
Epoch 00036: val_loss did not improve from 0.80483
45/45 [=====] - 21s 467ms/step - loss: 0.5338 - categorical_accuracy: 0.8370 - val_loss: 0.8935 - val_categorical_accuracy: 0.6714 - lr: 1.0000e-04
Epoch 37/50
45/45 [=====] - ETA: 0s - loss: 0.6265 - categorical_accuracy: 0.7852
Epoch 00037: val_loss did not improve from 0.80483
45/45 [=====] - 21s 474ms/step - loss: 0.6265 - categorical_accuracy: 0.7852 - val_loss: 0.8610 - val_categorical_accuracy: 0.7286 - lr: 1.0000e-04
Epoch 38/50
45/45 [=====] - ETA: 0s - loss: 0.5749 - categorical_accuracy: 0.8519
Epoch 00038: val_loss did not improve from 0.80483
45/45 [=====] - 21s 471ms/step - loss: 0.5749 - categorical_accuracy: 0.8519 - val_loss: 0.8346 - val_categorical_accuracy: 0.6286 - lr: 1.0000e-04
Epoch 39/50
45/45 [=====] - ETA: 0s - loss: 0.6854 - categorical_accuracy: 0.7926
Epoch 00039: val_loss did not improve from 0.80483
45/45 [=====] - 21s 472ms/step - loss: 0.6854 - categorical_accuracy: 0.7926 - val_loss: 1.0307 - val_categorical_accuracy: 0.5857 - lr: 1.0000e-05
Epoch 40/50
45/45 [=====] - ETA: 0s - loss: 0.6554 - categorical_accuracy: 0.7481
Epoch 00040: val_loss did not improve from 0.80483
45/45 [=====] - 21s 473ms/step - loss: 0.6554 - categorical_accuracy: 0.7481 - val_loss: 0.9844 - val_categorical_accuracy: 0.6000 - lr: 1.0000e-05
Epoch 42/50
45/45 [=====] - ETA: 0s - loss: 0.6525 - categorical_accuracy: 0.7704
Epoch 00042: val_loss did not improve from 0.80483
45/45 [=====] - 21s 478ms/step - loss: 0.6525 - categorical_accuracy: 0.7704 - val_loss: 1.0740 - val_categorical_accuracy: 0.6286 - lr: 1.0000e-05
Epoch 43/50
45/45 [=====] - ETA: 0s - loss: 0.5451 - categorical_accuracy: 0.8296
Epoch 00043: val_loss did not improve from 0.80483
45/45 [=====] - 21s 476ms/step - loss: 0.5451 - categorical_accuracy: 0.8296 - val_loss: 0.8655 - val_categorical_accuracy: 0.6857 - lr: 1.0000e-05

```

1 # Visualizing the Model Results using TensorBoard

2 %tensorboard --logdir logs

→ Reusing TensorBoard on port 6006 (pid 646), started 0:18:12 ago. (Use '!kill 646' to kill it.)



Best Epoch(Epoch 50) Result from the above Model1 **Train Accuracy : 76.30 Validation Accuracy : 74.29**

- We can clearly see that Model1 is performing better in terms of generalizing the model.
- But clearly the Accuracy is less on both the Train and Validation data.

CONV3D : Model 2

Conv 3D Model1 : (Batch Size: 10 , No. of Frames = 30 , Optimizer=Adam , Dense Layer= 64)

```

1 #Changing the batch Size
2 train_generator,val_generator,steps_per_epoch,validation_steps = gen_train_val(batch_size=20)
3
4 #Setting the Callback Parameters
5 callbacks_list2 = set_checkpoint('model2',pati=6)

1 # Model Architecture
2 model2 = Sequential([
3     Conv3D(16, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform', input_shape=(30,120,120,3)),
4     BatchNormalization(),
5     Conv3D(32, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
6     BatchNormalization(),
7     Conv3D(64, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
8     BatchNormalization(),
9     MaxPooling3D(pool_size=(2, 2, 2)),
10    Conv3D(128, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
11    BatchNormalization(),
12    MaxPooling3D(pool_size=(2, 2, 2)),
13    Flatten(),
14    Dense(64, activation='relu', kernel_initializer='he_uniform'),
15    BatchNormalization(),
16    Dense(5, activation='softmax')
17 ])

```

```

1 #Compiling the Model
2 optimiser = tf.keras.optimizers.Adam(learning_rate=0.001) #write your optimizer
3 model2.compile(optimizer=optimiser, loss='categorical_crossentropy', metrics=['categorical_accuracy'])
4 print (model2.summary())

```

→ Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv3d_10 (Conv3D)	(None, 28, 118, 118, 16)	1312
batch_normalization_10 (Batch Normalization)	(None, 28, 118, 118, 16)	64
conv3d_11 (Conv3D)	(None, 26, 116, 116, 32)	13856
batch_normalization_11 (Batch Normalization)	(None, 26, 116, 116, 32)	128
conv3d_12 (Conv3D)	(None, 24, 114, 114, 64)	55360
batch_normalization_12 (Batch Normalization)	(None, 24, 114, 114, 64)	256
max_pooling3d_6 (MaxPooling3D)	(None, 12, 57, 57, 64)	0
conv3d_13 (Conv3D)	(None, 10, 55, 55, 128)	221312
batch_normalization_13 (Batch Normalization)	(None, 10, 55, 55, 128)	512
max_pooling3d_7 (MaxPooling3D)	(None, 5, 27, 27, 128)	0
flatten_3 (Flatten)	(None, 466560)	0
dense_5 (Dense)	(None, 64)	29859904
batch_normalization_14 (Batch Normalization)	(None, 64)	256
dense_6 (Dense)	(None, 5)	325

Total params: 30,153,285
Trainable params: 30,152,677
Non-trainable params: 608

None

```

1 # Model Fit
2 model2.fit(train_generator,
3             steps_per_epoch=steps_per_epoch,
4             epochs=50,
5             verbose=1,
6             validation_data=val_generator,
7             validation_steps=validation_steps,
8             callbacks=callbacks_list2)

```

→ Source path = /content/train ; batch size = 20

Epoch 1/50
34/34 [=====] - ETA: 0s - loss: 1.3981 - categorical_accuracy: 0.4284
Source path = /content/val ; batch size = 20
Epoch 00001: val_loss improved from inf to 19.28531, saving model to model/model2.h5
34/34 [=====] - 61s 2s/step - loss: 1.3981 - categorical_accuracy: 0.4284 - val_loss: 19.2853 - val_categorical_accuracy: 0.2000 - lr: 0.0010
Epoch 2/50
34/34 [=====] - ETA: 0s - loss: 1.2044 - categorical_accuracy: 0.5098
Epoch 00002: val_loss improved from 19.28531 to 18.41357, saving model to model/model2.h5
34/34 [=====] - 27s 820ms/step - loss: 1.2044 - categorical_accuracy: 0.5098 - val_loss: 18.4136 - val_categorical_accuracy: 0.1800 - lr: 0.0010
Epoch 3/50
34/34 [=====] - ETA: 0s - loss: 1.3247 - categorical_accuracy: 0.4216
Epoch 00003: val_loss did not improve from 18.41357
34/34 [=====] - 25s 757ms/step - loss: 1.3247 - categorical_accuracy: 0.4216 - val_loss: 31.1670 - val_categorical_accuracy: 0.1900 - lr: 0.0010
Epoch 4/50
34/34 [=====] - ETA: 0s - loss: 1.4603 - categorical_accuracy: 0.3529
Epoch 00004: val_loss improved from 18.41357 to 3.05344, saving model to model/model2.h5
34/34 [=====] - 27s 800ms/step - loss: 1.4603 - categorical_accuracy: 0.3529 - val_loss: 3.0534 - val_categorical_accuracy: 0.3800 - lr: 0.0010
Epoch 5/50
34/34 [=====] - ETA: 0s - loss: 1.3089 - categorical_accuracy: 0.4608
Epoch 00005: val_loss did not improve from 3.05344
34/34 [=====] - 26s 771ms/step - loss: 1.3089 - categorical_accuracy: 0.4608 - val_loss: 3.9947 - val_categorical_accuracy: 0.3700 - lr: 0.0010
Epoch 6/50
34/34 [=====] - ETA: 0s - loss: 1.7235 - categorical_accuracy: 0.2353
Epoch 00006: val_loss did not improve from 3.05344
34/34 [=====] - 26s 775ms/step - loss: 1.7235 - categorical_accuracy: 0.2353 - val_loss: 10.2852 - val_categorical_accuracy: 0.1900 - lr: 0.0010
Epoch 7/50
34/34 [=====] - ETA: 0s - loss: 1.6958 - categorical_accuracy: 0.3725
Epoch 00007: val_loss did not improve from 3.05344
34/34 [=====] - 26s 766ms/step - loss: 1.6958 - categorical_accuracy: 0.3725 - val_loss: 4.2388 - val_categorical_accuracy: 0.2400 - lr: 0.0010
Epoch 8/50
34/34 [=====] - ETA: 0s - loss: 1.5186 - categorical_accuracy: 0.3529
Epoch 00008: val_loss improved from 3.05344 to 2.49613, saving model to model/model2.h5
34/34 [=====] - 26s 780ms/step - loss: 1.5186 - categorical_accuracy: 0.3529 - val_loss: 2.4961 - val_categorical_accuracy: 0.2500 - lr: 0.0010
Epoch 9/50
34/34 [=====] - ETA: 0s - loss: 1.5864 - categorical_accuracy: 0.3039
Epoch 00009: val_loss improved from 2.49613 to 2.23198, saving model to model/model2.h5
34/34 [=====] - 26s 790ms/step - loss: 1.5864 - categorical_accuracy: 0.3039 - val_loss: 2.2320 - val_categorical_accuracy: 0.3300 - lr: 0.0010
Epoch 10/50
34/34 [=====] - ETA: 0s - loss: 1.4191 - categorical_accuracy: 0.3235
Epoch 0010: val_loss improved from 2.23198 to 1.54431, saving model to model/model2.h5
34/34 [=====] - 26s 793ms/step - loss: 1.4191 - categorical_accuracy: 0.3235 - val_loss: 1.5443 - val_categorical_accuracy: 0.4900 - lr: 0.0010
Epoch 11/50
34/34 [=====] - ETA: 0s - loss: 1.6019 - categorical_accuracy: 0.3725
Epoch 0011: val_loss improved from 1.54431 to 1.41477, saving model to model/model2.h5
34/34 [=====] - 26s 790ms/step - loss: 1.6019 - categorical_accuracy: 0.3725 - val_loss: 1.4148 - val_categorical_accuracy: 0.4800 - lr: 0.0010
Epoch 12/50
34/34 [=====] - ETA: 0s - loss: 1.3308 - categorical_accuracy: 0.4706
Epoch 0012: val_loss improved from 1.41477 to 1.30888, saving model to model/model2.h5
34/34 [=====] - 26s 773ms/step - loss: 1.3308 - categorical_accuracy: 0.4706 - val_loss: 1.3089 - val_categorical_accuracy: 0.5100 - lr: 0.0010
Epoch 13/50
34/34 [=====] - ETA: 0s - loss: 1.3525 - categorical_accuracy: 0.4314
Epoch 0013: val_loss improved from 1.30888 to 1.28501, saving model to model/model2.h5
34/34 [=====] - 26s 791ms/step - loss: 1.3525 - categorical_accuracy: 0.4314 - val_loss: 1.2850 - val_categorical_accuracy: 0.5300 - lr: 0.0010
Epoch 14/50
34/34 [=====] - ETA: 0s - loss: 1.3801 - categorical_accuracy: 0.4412
Epoch 0014: val_loss improved from 1.28501 to 1.27724, saving model to model/model2.h5
34/34 [=====] - 26s 785ms/step - loss: 1.3801 - categorical_accuracy: 0.4412 - val_loss: 1.2772 - val_categorical_accuracy: 0.5500 - lr: 0.0010

Best Epoch (Epoch 34) Result from the above Model2: Train Accuracy : 69 Validation Accuracy : 64

- We can clearly see that Model1 is performing better in terms of generalizing the model.
- Accuracy is also less than Model 1

```

1 # Tenser Board
2 %tensorboard --logdir logs

```

→ Reusing TensorBoard on port 6006 (pid 646), started 0:48:56 ago. (Use '!kill 646' to kill it.)

TensorBoard

SCALARS GRAPHS DISTRIBUTIONS HISTOGRAMS TIME SERIES

INACTIVE

- Show data download links
- Ignore outliers in chart scaling

Tooltip sorting method: default ▾

Smoothing

0.6

Horizontal Axis

STEP RELATIVE WALL

Runs

Write a regex to filter runs

- model/train
- model/validation
- model1/train
- model1/validation
- model2/train
- model2/validation
- model3/train

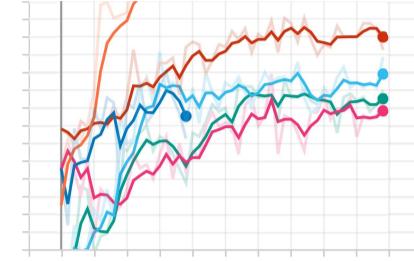
TOGGLE ALL RUNS

logs

Filter tags (regular expressions supported)

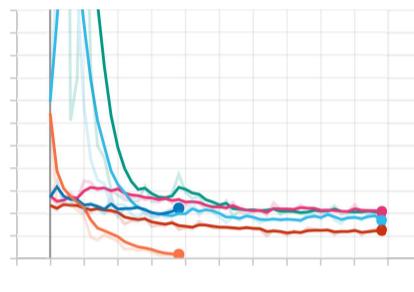
epoch_categorical_accuracy

epoch_categorical_accuracy
tag: epoch_categorical_accuracy



epoch_loss

epoch_loss
tag: epoch_loss



CONV3D : Model 3

(Batch Size: 15 , No. of Frames = 30 , Optimizer=Adam , LR =0.01,Dense_Layer=32)

```

1 #Changing the batch Size
2 train_generator,val_generator,steps_per_epoch,validation_steps = gen_train_val(batch_size=15)
3
4 #Setting the Callback Parameters
5 callbacks_list3 = set_checkpoint('model3',pati=6)

1 # Model Architecture
2 model3 = Sequential([
3     Conv3D(16, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform', input_shape=(30,120,120,3)),
4     BatchNormalization(),
5     Conv3D(32, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
6     BatchNormalization(),
7     Conv3D(64, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
8     BatchNormalization(),
9     MaxPooling3D(pool_size=(2, 2, 2)),
10    Conv3D(128, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform'),
11    BatchNormalization(),
12    MaxPooling3D(pool_size=(2, 2, 2)),
13    Flatten(),
14    Dense(32, activation='relu', kernel_initializer='he_uniform'),
15    #Dropout(0.1),
16    BatchNormalization(),
17    Dense(5, activation='softmax')
18 ])

1 ## Compiling the Model
2 optimiser = tf.keras.optimizers.Adam(learning_rate=0.01) # write your optimizer
3 model3.compile(optimizer=optimiser, loss='categorical_crossentropy', metrics=['categorical_accuracy'])
4 print (model3.summary())

```

→ Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
conv3d_14 (Conv3D)	(None, 28, 118, 118, 16)	1312
=====		
batch_normalization_15 (BatchNormalization)	(None, 28, 118, 118, 16)	64
conv3d_15 (Conv3D)	(None, 26, 116, 116, 32)	13856
batch_normalization_16 (BatchNormalization)	(None, 26, 116, 116, 32)	128
conv3d_16 (Conv3D)	(None, 24, 114, 114, 64)	55360
batch_normalization_17 (BatchNormalization)	(None, 24, 114, 114, 64)	256
max_pooling3d_8 (MaxPooling3D)	(None, 12, 57, 57, 64)	0
conv3d_17 (Conv3D)	(None, 10, 55, 55, 128)	221312
batch_normalization_18 (BatchNormalization)	(None, 10, 55, 55, 128)	512
max_pooling3d_9 (MaxPooling3D)	(None, 5, 27, 27, 128)	0
flatten_4 (Flatten)	(None, 466560)	0
dense_7 (Dense)	(None, 32)	14929952

```
batch_normalization_19 (Batch Normalization)          128
dense_8 (Dense)                                     (None, 5)           165
=====
Total params: 15,223,045
Trainable params: 15,222,501
Non-trainable params: 544
```

None

```
1 ## Model Fit
2 model3.fit(train_generator,
3             steps_per_epoch=steps_per_epoch,
4             epochs=50,
5             verbose=1,
6             validation_data=val_generator,
7             validation_steps=validation_steps,
8             callbacks=callbacks_list3)

→ Source path = /content/train ; batch size = 15
Epoch 1/50
45/45 [=====] - ETA: 0s - loss: 1.5155 - categorical_accuracy: 0.3680Source path = /content/val ; batch size = 15

Epoch 00001: val_loss improved from inf to 119.81718, saving model to model/model3.h5
45/45 [=====] - 53s 1s/step - loss: 1.5155 - categorical_accuracy: 0.3680 - val_loss: 119.8172 - val_categorical_accuracy: 0.2400 - lr: 0.0100
Epoch 2/50
45/45 [=====] - ETA: 0s - loss: 1.5756 - categorical_accuracy: 0.3333
Epoch 00002: val_loss improved from 119.81718 to 47.96495, saving model to model/model3.h5
45/45 [=====] - 22s 501ms/step - loss: 1.5756 - categorical_accuracy: 0.3333 - val_loss: 47.9649 - val_categorical_accuracy: 0.2000 - lr: 0.0100
Epoch 3/50
45/45 [=====] - ETA: 0s - loss: 1.5092 - categorical_accuracy: 0.3481
Epoch 00003: val_loss improved from 47.96495 to 10.47097, saving model to model/model3.h5
45/45 [=====] - 22s 499ms/step - loss: 1.5092 - categorical_accuracy: 0.3481 - val_loss: 10.4710 - val_categorical_accuracy: 0.2286 - lr: 0.0100
Epoch 4/50
45/45 [=====] - ETA: 0s - loss: 1.5566 - categorical_accuracy: 0.3185
Epoch 00004: val_loss did not improve from 10.47097
45/45 [=====] - 22s 504ms/step - loss: 1.5566 - categorical_accuracy: 0.3185 - val_loss: 12.9208 - val_categorical_accuracy: 0.2143 - lr: 0.0100
Epoch 5/50
45/45 [=====] - ETA: 0s - loss: 1.5159 - categorical_accuracy: 0.3852
Epoch 00005: val_loss improved from 10.47097 to 1.71418, saving model to model/model3.h5
45/45 [=====] - 22s 499ms/step - loss: 1.5159 - categorical_accuracy: 0.3852 - val_loss: 1.7142 - val_categorical_accuracy: 0.3571 - lr: 0.0100
Epoch 6/50
45/45 [=====] - ETA: 0s - loss: 1.4723 - categorical_accuracy: 0.3407
Epoch 00006: val_loss did not improve from 1.71418
45/45 [=====] - 22s 495ms/step - loss: 1.4723 - categorical_accuracy: 0.3407 - val_loss: 2.3845 - val_categorical_accuracy: 0.3286 - lr: 0.0100
Epoch 7/50
45/45 [=====] - ETA: 0s - loss: 1.4102 - categorical_accuracy: 0.3704
Epoch 00007: val_loss did not improve from 1.71418
45/45 [=====] - 22s 499ms/step - loss: 1.4102 - categorical_accuracy: 0.3704 - val_loss: 2.6589 - val_categorical_accuracy: 0.3429 - lr: 0.0100
Epoch 8/50
45/45 [=====] - ETA: 0s - loss: 1.4636 - categorical_accuracy: 0.3037
Epoch 00008: val_loss did not improve from 1.71418
45/45 [=====] - 22s 494ms/step - loss: 1.4636 - categorical_accuracy: 0.3037 - val_loss: 2.9771 - val_categorical_accuracy: 0.2429 - lr: 0.0100
Epoch 9/50
45/45 [=====] - ETA: 0s - loss: 1.4042 - categorical_accuracy: 0.4074
Epoch 00009: val_loss improved from 1.71418 to 1.13578, saving model to model/model3.h5
45/45 [=====] - 22s 494ms/step - loss: 1.4042 - categorical_accuracy: 0.4074 - val_loss: 1.1358 - val_categorical_accuracy: 0.5000 - lr: 0.0100
Epoch 10/50
45/45 [=====] - ETA: 0s - loss: 1.4263 - categorical_accuracy: 0.3778
Epoch 0010: val_loss did not improve from 1.13578
45/45 [=====] - 22s 494ms/step - loss: 1.4263 - categorical_accuracy: 0.3778 - val_loss: 1.1855 - val_categorical_accuracy: 0.4429 - lr: 0.0100
Epoch 11/50
45/45 [=====] - ETA: 0s - loss: 1.2764 - categorical_accuracy: 0.4222
Epoch 0011: val_loss improved from 1.13578 to 1.04099, saving model to model/model3.h5
45/45 [=====] - 22s 503ms/step - loss: 1.2764 - categorical_accuracy: 0.4222 - val_loss: 1.0410 - val_categorical_accuracy: 0.5714 - lr: 0.0100
Epoch 12/50
45/45 [=====] - ETA: 0s - loss: 1.2662 - categorical_accuracy: 0.4667
Epoch 0012: val_loss did not improve from 1.04099
45/45 [=====] - 22s 490ms/step - loss: 1.2662 - categorical_accuracy: 0.4667 - val_loss: 3.2940 - val_categorical_accuracy: 0.3143 - lr: 0.0100
Epoch 13/50
45/45 [=====] - ETA: 0s - loss: 1.4883 - categorical_accuracy: 0.3556
Epoch 0013: val_loss did not improve from 1.04099
45/45 [=====] - 22s 483ms/step - loss: 1.4883 - categorical_accuracy: 0.3556 - val_loss: 57.5957 - val_categorical_accuracy: 0.1857 - lr: 0.0100
Epoch 14/50
45/45 [=====] - ETA: 0s - loss: 1.4666 - categorical_accuracy: 0.3704
Epoch 0014: val_loss did not improve from 1.04099
45/45 [=====] - 22s 487ms/step - loss: 1.4666 - categorical_accuracy: 0.3704 - val_loss: 5.7077 - val_categorical_accuracy: 0.3000 - lr: 0.0100

1 # Visualizing the Model Results using TensorBoard
2
3 %tensorboard --logdir logs
```

TensorBoard

SCALARS GRAPHS DISTRIBUTIONS HISTOGRAMS TIME SERIES

INACTIVE

- Show data download links
- Ignore outliers in chart scaling

Tooltip sorting method: default ▾

Smoothing

 0.6

Horizontal Axis

STEP RELATIVE WALL

Runs

Write a regex to filter runs

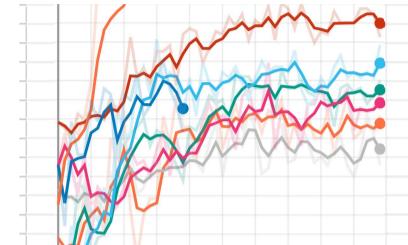
- model/train
- model/validation
- model1/train
- model1/validation
- model2/train
- model2/validation
- model3/train
- model3/validation

TOGGLE ALL RUNS

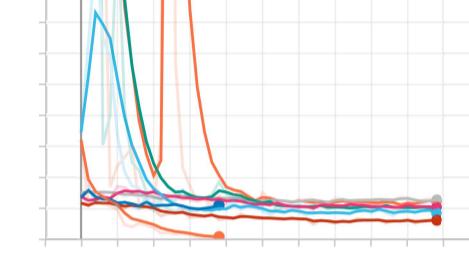
logs

Filter tags (regular expressions supported)

epoch_categorical_accuracy

epoch_categorical_accuracy
tag: epoch_categorical_accuracy

epoch_loss

epoch_loss
tag: epoch_loss

Best Epoch (Epoch 48) Result from the above Model3: Train Accuracy : 56 Validation Accuracy : 54

Conv3D: Model 4

- (Batch Size: 7 , No. of Frames = 30 , Optimizer=Adam , LR =0.0001,Dense_Layer=32)

```

1 #Changing the batch Size
2 train_generator,val_generator,steps_per_epoch,validation_steps = gen_train_val(batch_size=7)
3
4 #Setting the Callback Parameters
5 callbacks_list3 = set_checkpoint('model4',pati=6)

1 model4 = Sequential([
2     Conv3D(16, kernel_size=(3, 3, 3), activation='relu', kernel_initializer='he_uniform',padding='same', input_shape=(30,120,120,3)),
3     BatchNormalization(),
4     #MaxPooling3D(pool_size=(2, 2, 2)),
5     Conv3D(32, kernel_size=(3, 3, 3), activation='relu',padding='same', kernel_initializer='he_uniform'),
6     BatchNormalization(),
7     MaxPooling3D(pool_size=(2, 2, 2)),
8     Conv3D(64, kernel_size=(3, 3, 3), activation='relu',padding='same', kernel_initializer='he_uniform'),
9     BatchNormalization(),
10    MaxPooling3D(pool_size=(2, 2, 2)),
11    Conv3D(128, kernel_size=(3, 3, 3), activation='relu',padding='same', kernel_initializer='he_uniform'),
12    BatchNormalization(),
13    MaxPooling3D(pool_size=(2, 2, 2)),
14    Flatten(),
15    Dense(32, activation='relu', kernel_initializer='he_uniform'),
16    #Dropout(0.25),
17    BatchNormalization(),
18    Dense(5, activation='softmax')
19 ])

1 optimiser = tf.keras.optimizers.Adam(learning_rate=0.0001) # write your optimizer
2 model4.compile(optimizer=optimiser, loss='categorical_crossentropy', metrics=['categorical_accuracy'])
3 print (model4.summary())

```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
=====		
conv3d_18 (Conv3D)	(None, 30, 120, 120, 16)	1312
=====		
batch_normalization_20 (BatchNormalization)	(None, 30, 120, 120, 16)	64
conv3d_19 (Conv3D)	(None, 30, 120, 120, 32)	13856
batch_normalization_21 (BatchNormalization)	(None, 30, 120, 120, 32)	128
max_pooling3d_10 (MaxPooling3D)	(None, 15, 60, 60, 32)	0
conv3d_20 (Conv3D)	(None, 15, 60, 60, 64)	55360
batch_normalization_22 (BatchNormalization)	(None, 15, 60, 60, 64)	256
max_pooling3d_11 (MaxPooling3D)	(None, 7, 30, 30, 64)	0
conv3d_21 (Conv3D)	(None, 7, 30, 30, 128)	221312
batch_normalization_23 (BatchNormalization)	(None, 7, 30, 30, 128)	512
max_pooling3d_12 (MaxPooling3D)	(None, 3, 15, 15, 128)	0
flatten_5 (Flatten)	(None, 86400)	0

```

dense_9 (Dense)          (None, 32)           2764832
batch_normalization_24 (Batch Normalization) (None, 32)           128
dense_10 (Dense)          (None, 5)            165
=====
Total params: 3,057,925
Trainable params: 3,057,381
Non-trainable params: 544

```

None

```

1 model4.fit(train_generator,
2     steps_per_epoch=steps_per_epoch,
3     epochs=50,
4     verbose=1,
5     validation_data=val_generator,
6     validation_steps=validation_steps,
7     callbacks=callbacks_list)

Source path = /content/train ; batch size = 7
Epoch 1/50
95/95 [=====] - ETA: 0s - loss: 1.3216 - categorical_accuracy: 0.4691Source path = /content/val ; batch size = 7

Epoch 00001: val_loss did not improve from 0.94098
95/95 [=====] - 49s 498ms/step - loss: 1.3216 - categorical_accuracy: 0.4691 - val_loss: 3.5595 - val_categorical_accuracy: 0.2100 - lr: 1.0000e-04
Epoch 2/50
95/95 [=====] - ETA: 0s - loss: 0.7929 - categorical_accuracy: 0.7305
Epoch 00002: val_loss did not improve from 0.94098
95/95 [=====] - 32s 337ms/step - loss: 0.7929 - categorical_accuracy: 0.7305 - val_loss: 3.0417 - val_categorical_accuracy: 0.2333 - lr: 1.0000e-04
Epoch 3/50
95/95 [=====] - ETA: 0s - loss: 0.7289 - categorical_accuracy: 0.7549
Epoch 00003: val_loss did not improve from 0.94098
95/95 [=====] - 26s 273ms/step - loss: 0.7289 - categorical_accuracy: 0.7549 - val_loss: 1.9516 - val_categorical_accuracy: 0.3667 - lr: 1.0000e-04
Epoch 4/50
95/95 [=====] - ETA: 0s - loss: 0.9127 - categorical_accuracy: 0.6561
Epoch 00004: val_loss did not improve from 0.94098
95/95 [=====] - 22s 234ms/step - loss: 0.9127 - categorical_accuracy: 0.6561 - val_loss: 1.8493 - val_categorical_accuracy: 0.4000 - lr: 1.0000e-04
Epoch 5/50
95/95 [=====] - ETA: 0s - loss: 0.9673 - categorical_accuracy: 0.6351
Epoch 00005: val_loss did not improve from 0.94098
95/95 [=====] - 21s 225ms/step - loss: 0.9673 - categorical_accuracy: 0.6351 - val_loss: 1.5497 - val_categorical_accuracy: 0.4667 - lr: 1.0000e-04
Epoch 6/50
95/95 [=====] - ETA: 0s - loss: 0.9127 - categorical_accuracy: 0.6842
Epoch 00006: val_loss did not improve from 0.94098
95/95 [=====] - 21s 225ms/step - loss: 0.9127 - categorical_accuracy: 0.6842 - val_loss: 1.0356 - val_categorical_accuracy: 0.5667 - lr: 1.0000e-04
Epoch 7/50
95/95 [=====] - ETA: 0s - loss: 0.8693 - categorical_accuracy: 0.6982
Epoch 00007: val_loss did not improve from 0.94098
95/95 [=====] - 21s 226ms/step - loss: 0.8693 - categorical_accuracy: 0.6982 - val_loss: 0.9585 - val_categorical_accuracy: 0.6333 - lr: 1.0000e-04
Epoch 8/50
95/95 [=====] - ETA: 0s - loss: 0.6846 - categorical_accuracy: 0.7895
Epoch 00008: val_loss did not improve from 0.94098
95/95 [=====] - 22s 230ms/step - loss: 0.6846 - categorical_accuracy: 0.7895 - val_loss: 0.9784 - val_categorical_accuracy: 0.5667 - lr: 1.0000e-04
Epoch 9/50
95/95 [=====] - ETA: 0s - loss: 0.7579 - categorical_accuracy: 0.7298
Epoch 00009: val_loss did not improve from 0.94098
95/95 [=====] - 22s 230ms/step - loss: 0.7579 - categorical_accuracy: 0.7298 - val_loss: 0.9447 - val_categorical_accuracy: 0.6000 - lr: 1.0000e-04
Epoch 10/50
95/95 [=====] - ETA: 0s - loss: 0.7182 - categorical_accuracy: 0.7754
Epoch 0010: val_loss improved from 0.94098 to 0.87866, saving model to model/model.h5
95/95 [=====] - 22s 231ms/step - loss: 0.7182 - categorical_accuracy: 0.7754 - val_loss: 0.8787 - val_categorical_accuracy: 0.7333 - lr: 1.0000e-04
Epoch 11/50
95/95 [=====] - ETA: 0s - loss: 0.5791 - categorical_accuracy: 0.8632
Epoch 0011: val_loss did not improve from 0.87866
95/95 [=====] - 21s 225ms/step - loss: 0.5791 - categorical_accuracy: 0.8632 - val_loss: 0.9839 - val_categorical_accuracy: 0.6667 - lr: 1.0000e-04
Epoch 12/50
95/95 [=====] - ETA: 0s - loss: 0.6253 - categorical_accuracy: 0.8211
Epoch 0012: val_loss did not improve from 0.87866
95/95 [=====] - 23s 240ms/step - loss: 0.6253 - categorical_accuracy: 0.8211 - val_loss: 1.5694 - val_categorical_accuracy: 0.4667 - lr: 1.0000e-04
Epoch 13/50
95/95 [=====] - ETA: 0s - loss: 0.6426 - categorical_accuracy: 0.8105
Epoch 0013: val_loss improved from 0.87866 to 0.76935, saving model to model/model.h5
95/95 [=====] - 22s 230ms/step - loss: 0.6426 - categorical_accuracy: 0.8105 - val_loss: 0.7694 - val_categorical_accuracy: 0.7333 - lr: 1.0000e-04
Epoch 14/50
95/95 [=====] - ETA: 0s - loss: 0.6050 - categorical_accuracy: 0.8000
Epoch 0014: val_loss did not improve from 0.76935
95/95 [=====] - 21s 223ms/step - loss: 0.6050 - categorical_accuracy: 0.8000 - val_loss: 0.8425 - val_categorical_accuracy: 0.6333 - lr: 1.0000e-04

```

```

1 # Visualizing the Model Results using TensorBoard
2
3 %tensorboard --logdir logs

```

TensorBoard

SCALARS GRAPHS DISTRIBUTIONS HISTOGRAMS TIME SERIES

INACTIVE

- Show data download links
- Ignore outliers in chart scaling

Tooltip sorting method: default ▾

Smoothing

0.6

Horizontal Axis

STEP RELATIVE WALL

Runs

Write a regex to filter runs

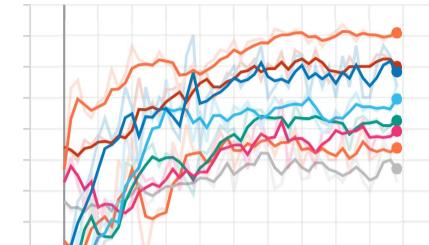
- model/train
- model/validation
- model1/train
- model1/validation
- model2/train
- model2/validation
- model3/train
- model3/validation
- CNN_RNN/train

TOGGLE ALL RUNS

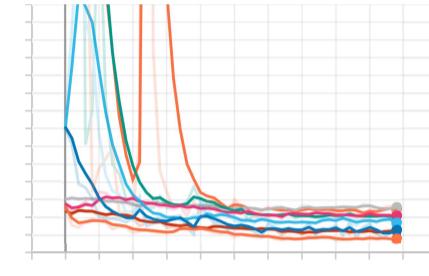
logs

Filter tags (regular expressions supported)

epoch_categorical_accuracy

epoch_categorical_accuracy
tag: epoch_categorical_accuracy

epoch_loss

epoch_loss
tag: epoch_loss

Best Epoch (Epoch 88) Result from the above Model3: Train Accuracy : 89 Validation Accuracy : 87

- We can clearly see that Model3 is performing better in terms of generalizing the model.
- Validation Accuracy is also also better than Model 1 and Model 2

Models : CNN + RNN

- Now we will be working on CNN + RNN Models and Transfer Learning Models

```

1 #Changing the batch Size
2 train_generator, val_generator, steps_per_epoch, validation_steps = gen_train_val(batch_size=7)
3
4 #Setting the Callback Parameters
5 callbacks_list5 = set_checkpoint('CNN_RNN',pati=6)

1 CNN_RNN = Sequential([
2     TimeDistributed(Conv2D(16,kernel_size=(3,3),activation='relu',padding='same',kernel_regularizer='l2',bias_initializer='he_uniform',bias_regularizer='l2')),input_shape=(30,120,120,3)
3     TimeDistributed(BatchNormalization()),
4     TimeDistributed(MaxPooling2D()),
5
6     TimeDistributed(Conv2D(32,kernel_size=(3,3),activation='relu',padding='same',kernel_regularizer='l2',bias_initializer='he_uniform',bias_regularizer='l2')),
7     TimeDistributed(BatchNormalization()),
8     # TimeDistributed(MaxPooling2D()),
9
10    TimeDistributed(Conv2D(64,kernel_size=(3,3),activation='relu',padding='same',kernel_regularizer='l2',bias_initializer='he_uniform',bias_regularizer='l2')),
11    TimeDistributed(BatchNormalization()),
12    # TimeDistributed(MaxPooling2D()),
13
14    TimeDistributed(Conv2D(128,kernel_size=(3,3),activation='relu',padding='same',kernel_regularizer='l2',bias_initializer='he_uniform',bias_regularizer='l2')),
15    TimeDistributed(BatchNormalization()),
16    TimeDistributed(MaxPooling2D()),
17
18    TimeDistributed(Conv2D(256,kernel_size=(3,3),activation='relu',padding='same',kernel_regularizer='l2',bias_initializer='he_uniform',bias_regularizer='l2')),
19    TimeDistributed(BatchNormalization()),
20    TimeDistributed(MaxPooling2D()),
21    TimeDistributed(Flatten()),
22
23    # TimeDistributed(Dropout(0.25)),
24    # TimeDistributed(Dense(64,activation='relu',kernel_regularizer='l2')),
25
26    GRU(16,return_sequences=False),
27
28    Dense(16,activation='relu',kernel_regularizer='l2'),
29
30    Dense(5, activation='softmax')
31 ])

```

```

1 optimiser = tf.keras.optimizers.SGD(learning_rate=0.01) # write your optimizer
2 CNN_RNN.compile(optimizer=optimiser, loss='categorical_crossentropy', metrics=['CategoricalAccuracy'])
3 CNN_RNN.summary()

```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
<hr/>		
time_distributed	(TimeDistr (None, 30, 120, 120, 16)	448 ibuted)
time_distributed_1	(TimeDis (None, 30, 120, 120, 16)	64 tributed)
time_distributed_2	(TimeDis (None, 30, 60, 60, 16)	0 tributed)
time_distributed_3	(TimeDis (None, 30, 60, 60, 32)	4640 tributed)
time_distributed_4	(TimeDis (None, 30, 60, 60, 32)	128

```
tributed)
```

```
time_distributed_5 (TimeDis (None, 30, 60, 60, 64) 18496
```

```
tributed)
```

```
time_distributed_6 (TimeDis (None, 30, 60, 60, 64) 256
```

```
tributed)
```

```
time_distributed_7 (TimeDis (None, 30, 60, 60, 128) 73856
```

```
tributed)
```

```
time_distributed_8 (TimeDis (None, 30, 60, 60, 128) 512
```

```
tributed)
```

```
time_distributed_9 (TimeDis (None, 30, 30, 30, 128) 0
```

```
tributed)
```

```
time_distributed_10 (TimeDi (None, 30, 30, 30, 256) 295168
```

```
stributed)
```

```
time_distributed_11 (TimeDi (None, 30, 30, 30, 256) 1024
```

```
stributed)
```

```
time_distributed_12 (TimeDi (None, 30, 15, 15, 256) 0
```

```
stributed)
```

```
time_distributed_13 (TimeDi (None, 30, 57600) 0
```

```
stributed)
```

```
gru (GRU) (None, 16) 2765664
```

```
dense_11 (Dense) (None, 16) 272
```

```
dense_12 (Dense) (None, 5) 85
```

```
=====
```

```
Total params: 3,160,613
```

```
Trainable params: 3,159,621
```

```
Non-trainable params: 992
```

```
1 CNN_RNN.fit(train_generator,
2     steps_per_epoch=steps_per_epoch,
3     epochs=50,
4     verbose=1,
5     validation_data=val_generator,
6     validation_steps=validation_steps,
7     callbacks=callbacks_list5)
8
```

Source path = /content/train ; batch size = 7
Epoch 1/50
95/95 [=====] - ETA: 0s - loss: 5.0923 - categorical_accuracy: 0.2081Source path = /content/val ; batch size = 7

Epoch 00001: val_loss improved from inf to 5.25212, saving model to model/CNN_RNN.h5
95/95 [=====] - 52s 503ms/step - loss: 5.0923 - categorical_accuracy: 0.2081 - val_loss: 5.2521 - val_categorical_accuracy: 0.1600 - lr: 0.0100
Epoch 2/50
95/95 [=====] - ETA: 0s - loss: 4.8982 - categorical_accuracy: 0.2821
Epoch 00002: val_loss improved from 5.25212 to 4.90046, saving model to model/CNN_RNN.h5
95/95 [=====] - 30s 315ms/step - loss: 4.8982 - categorical_accuracy: 0.2821 - val_loss: 4.9005 - val_categorical_accuracy: 0.2000 - lr: 0.0100
Epoch 3/50
95/95 [=====] - ETA: 0s - loss: 4.7272 - categorical_accuracy: 0.3287
Epoch 00003: val_loss improved from 4.90046 to 4.79702, saving model to model/CNN_RNN.h5
95/95 [=====] - 23s 239ms/step - loss: 4.7272 - categorical_accuracy: 0.3287 - val_loss: 4.7970 - val_categorical_accuracy: 0.1667 - lr: 0.0100
Epoch 4/50
95/95 [=====] - ETA: 0s - loss: 4.6158 - categorical_accuracy: 0.2912
Epoch 00004: val_loss improved from 4.79702 to 4.62067, saving model to model/CNN_RNN.h5
95/95 [=====] - 19s 199ms/step - loss: 4.6158 - categorical_accuracy: 0.2912 - val_loss: 4.6207 - val_categorical_accuracy: 0.2000 - lr: 0.0100
Epoch 5/50
95/95 [=====] - ETA: 0s - loss: 4.4826 - categorical_accuracy: 0.3509
Epoch 00005: val_loss improved from 4.62067 to 4.35728, saving model to model/CNN_RNN.h5
95/95 [=====] - 19s 206ms/step - loss: 4.4826 - categorical_accuracy: 0.3509 - val_loss: 4.3573 - val_categorical_accuracy: 0.3667 - lr: 0.0100
Epoch 6/50
95/95 [=====] - ETA: 0s - loss: 4.2950 - categorical_accuracy: 0.4105
Epoch 00006: val_loss improved from 4.35728 to 4.11698, saving model to model/CNN_RNN.h5
95/95 [=====] - 19s 203ms/step - loss: 4.2950 - categorical_accuracy: 0.4105 - val_loss: 4.1170 - val_categorical_accuracy: 0.5000 - lr: 0.0100
Epoch 7/50
95/95 [=====] - ETA: 0s - loss: 4.2201 - categorical_accuracy: 0.3298
Epoch 00007: val_loss did not improve from 4.11698
95/95 [=====] - 19s 198ms/step - loss: 4.2201 - categorical_accuracy: 0.3298 - val_loss: 4.1854 - val_categorical_accuracy: 0.3333 - lr: 0.0100
Epoch 8/50
95/95 [=====] - ETA: 0s - loss: 4.0627 - categorical_accuracy: 0.4070
Epoch 00008: val_loss improved from 4.11698 to 4.06504, saving model to model/CNN_RNN.h5
95/95 [=====] - 18s 195ms/step - loss: 4.0627 - categorical_accuracy: 0.4070 - val_loss: 4.0650 - val_categorical_accuracy: 0.3333 - lr: 0.0100
Epoch 9/50
95/95 [=====] - ETA: 0s - loss: 3.9726 - categorical_accuracy: 0.3930
Epoch 00009: val_loss did not improve from 4.06504
95/95 [=====] - 19s 202ms/step - loss: 3.9726 - categorical_accuracy: 0.3930 - val_loss: 4.1699 - val_categorical_accuracy: 0.2667 - lr: 0.0100
Epoch 10/50
95/95 [=====] - ETA: 0s - loss: 3.9525 - categorical_accuracy: 0.3579
Epoch 00010: val_loss improved from 4.06504 to 3.66747, saving model to model/CNN_RNN.h5
95/95 [=====] - 19s 206ms/step - loss: 3.9525 - categorical_accuracy: 0.3579 - val_loss: 3.6675 - val_categorical_accuracy: 0.5000 - lr: 0.0100
Epoch 11/50
95/95 [=====] - ETA: 0s - loss: 3.8151 - categorical_accuracy: 0.4246
Epoch 00011: val_loss did not improve from 3.66747
95/95 [=====] - 19s 196ms/step - loss: 3.8151 - categorical_accuracy: 0.4246 - val_loss: 3.7786 - val_categorical_accuracy: 0.3333 - lr: 0.0100
Epoch 12/50
95/95 [=====] - ETA: 0s - loss: 3.6483 - categorical_accuracy: 0.4211
Epoch 00012: val_loss did not improve from 3.66747
95/95 [=====] - 19s 199ms/step - loss: 3.6483 - categorical_accuracy: 0.4211 - val_loss: 3.7847 - val_categorical_accuracy: 0.3667 - lr: 0.0100
Epoch 13/50
95/95 [=====] - ETA: 0s - loss: 3.6275 - categorical_accuracy: 0.3895
Epoch 00013: val_loss improved from 3.66747 to 3.53565, saving model to model/CNN_RNN.h5
95/95 [=====] - 20s 211ms/step - loss: 3.6275 - categorical_accuracy: 0.3895 - val_loss: 3.5356 - val_categorical_accuracy: 0.4333 - lr: 0.0100
Epoch 14/50
95/95 [=====] - ETA: 0s - loss: 3.6418 - categorical_accuracy: 0.3895
Epoch 00014: val_loss did not improve from 3.53565
95/95 [=====] - 18s 190ms/step - loss: 3.6418 - categorical_accuracy: 0.3895 - val_loss: 3.7158 - val_categorical_accuracy: 0.3333 - lr: 0.0100

```
1 # Visualizing the Model Results using TenserBoard
```

```
2
```

```
3 %tensorboard --logdir logs
```

TensorBoard

SCALARS GRAPHS DISTRIBUTIONS HISTOGRAMS TIME SERIES

INACTIVE

- Show data download links
- Ignore outliers in chart scaling

Tooltip sorting method: default ▾

Smoothing

 0.6

Horizontal Axis

STEP RELATIVE WALL

Runs

Write a regex to filter runs

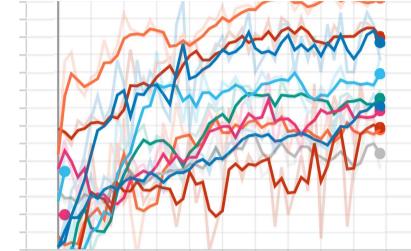
- model/train
- model/validation
- model1/train
- model1/validation
- model2/train
- model2/validation
- model3/train
- model3/validation
- CNN_RNN/train

TOGGLE ALL RUNS

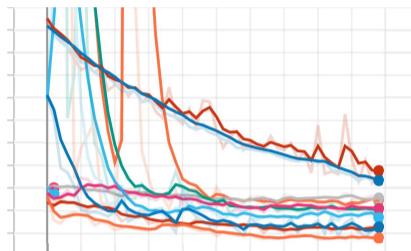
logs

Filter tags (regular expressions supported)

epoch_categorical_accuracy

epoch_categorical_accuracy
tag: epoch_categorical_accuracy

epoch_loss

epoch_loss
tag: epoch_loss

Best Epoch Result from the above Model: Train Accuracy : 63 /Validation Accuracy : 57

Transfer Learning Model (Mobile Net)

- Creating a Model using transfer Learning , we will be using MobileNet architecture for Transfer learning
- Transfer Learning has less parameters

```

1 #Changing the batch Size
2 train_generator,val_generator,steps_per_epoch,validation_steps = gen_train_val(batch_size=7)
3
4 #Setting the Callback Parameters
5 callbacks_list6 = set_checkpoint('MobileNet',pati=6)

```

```

1 ## Transfer Learning using Mobile Net Architecture
2 mobilenet = tf.keras.applications.MobileNetV2(include_top=False)
3 mobilenet.trainable = True
4 mobilnet = Sequential([
5     TimeDistributed(mobilenet,input_shape=(30,120,120,3)),
6     TimeDistributed(BatchNormalization()),
7     TimeDistributed(Flatten()),
8     GRU(6,return_sequences=False,dropout=0.25),
9     #BatchNormalization(),
10    Dense(5,activation='softmax')
11 ])

```

WARNING:tensorflow:`input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.
 Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5
 9412608/9406464 [=====] - 0s 0us/step
 9420800/9406464 [=====] - 0s 0us/step

```

1 ## Compile the Model
2 optimiser = tf.keras.optimizers.SGD(learning_rate=0.001) # write your optimizer
3 mobilnet.compile(optimizer=optimiser, loss='categorical_crossentropy', metrics=['CategoricalAccuracy'])
4 mobilnet.summary()

```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
<hr/>		
time_distributed_14 (TimeDistributed)	(None, 30, 3, 3, 1280)	2257984
time_distributed_15 (TimeDistributed)	(None, 30, 3, 3, 1280)	5120
time_distributed_16 (TimeDistributed)	(None, 30, 11520)	0
gru_1 (GRU)	(None, 6)	207504
dense_13 (Dense)	(None, 5)	35
<hr/>		
Total params:	2,470,643	
Trainable params:	2,433,971	
Non-trainable params:	36,672	

```

1 #Transfer Learning Model fit
2 mobilnet.fit(train_generator,
3               steps_per_epoch=steps_per_epoch,
4               epochs=50,
5               verbose=1,
6               callbacks=callbacks_list6,
7               validation_data=val_generator,
8               validation_steps=validation_steps)

```

```

Source path = /content/train ; batch size = 7
Epoch 1/50
95/95 [=====] - ETA: 0s - loss: 1.5286 - categorical_accuracy: 0.3379Source path = /content/val ; batch size = 7

Epoch 00001: val_loss improved from inf to 1.48336, saving model to model/MobileNet.h5
95/95 [=====] - 57s 537ms/step - loss: 1.5286 - categorical_accuracy: 0.3379 - val_loss: 1.4834 - val_categorical_accuracy: 0.3000 - lr: 0.0010
Epoch 2/50
95/95 [=====] - ETA: 0s - loss: 1.2885 - categorical_accuracy: 0.4716
Epoch 00002: val_loss did not improve from 1.48336
95/95 [=====] - 33s 354ms/step - loss: 1.2885 - categorical_accuracy: 0.4716 - val_loss: 1.5226 - val_categorical_accuracy: 0.3000 - lr: 0.0010
Epoch 3/50
95/95 [=====] - ETA: 0s - loss: 1.2455 - categorical_accuracy: 0.5014
Epoch 00003: val_loss improved from 1.48336 to 1.13933, saving model to model/MobileNet.h5
95/95 [=====] - 29s 302ms/step - loss: 1.2455 - categorical_accuracy: 0.5014 - val_loss: 1.1393 - val_categorical_accuracy: 0.5333 - lr: 0.0010
Epoch 4/50
95/95 [=====] - ETA: 0s - loss: 1.2300 - categorical_accuracy: 0.5018
Epoch 00004: val_loss did not improve from 1.13933
95/95 [=====] - 23s 244ms/step - loss: 1.2300 - categorical_accuracy: 0.5018 - val_loss: 1.1589 - val_categorical_accuracy: 0.5333 - lr: 0.0010
Epoch 5/50
95/95 [=====] - ETA: 0s - loss: 1.1705 - categorical_accuracy: 0.5719
Epoch 00005: val_loss improved from 1.13933 to 1.09619, saving model to model/MobileNet.h5
95/95 [=====] - 24s 255ms/step - loss: 1.1705 - categorical_accuracy: 0.5719 - val_loss: 1.0962 - val_categorical_accuracy: 0.6000 - lr: 0.0010
Epoch 6/50
95/95 [=====] - ETA: 0s - loss: 1.1863 - categorical_accuracy: 0.5123
Epoch 00006: val_loss improved from 1.09619 to 0.89137, saving model to model/MobileNet.h5
95/95 [=====] - 23s 245ms/step - loss: 1.1863 - categorical_accuracy: 0.5123 - val_loss: 0.8914 - val_categorical_accuracy: 0.7000 - lr: 0.0010
Epoch 7/50
95/95 [=====] - ETA: 0s - loss: 1.0956 - categorical_accuracy: 0.6105
Epoch 00007: val_loss did not improve from 0.89137
95/95 [=====] - 24s 249ms/step - loss: 1.0956 - categorical_accuracy: 0.6105 - val_loss: 1.1298 - val_categorical_accuracy: 0.5667 - lr: 0.0010
Epoch 8/50
95/95 [=====] - ETA: 0s - loss: 1.0659 - categorical_accuracy: 0.6140
Epoch 00008: val_loss did not improve from 0.89137
95/95 [=====] - 23s 239ms/step - loss: 1.0659 - categorical_accuracy: 0.6140 - val_loss: 1.1294 - val_categorical_accuracy: 0.5000 - lr: 0.0010
Epoch 9/50
95/95 [=====] - ETA: 0s - loss: 1.0268 - categorical_accuracy: 0.6175
Epoch 00009: val_loss did not improve from 0.89137
95/95 [=====] - 25s 261ms/step - loss: 1.0268 - categorical_accuracy: 0.6175 - val_loss: 1.2509 - val_categorical_accuracy: 0.5000 - lr: 0.0010
Epoch 10/50
95/95 [=====] - ETA: 0s - loss: 0.9709 - categorical_accuracy: 0.6491
Epoch 0010: val_loss improved from 0.89137 to 0.85638, saving model to model/MobileNet.h5
95/95 [=====] - 24s 251ms/step - loss: 0.9709 - categorical_accuracy: 0.6491 - val_loss: 0.8564 - val_categorical_accuracy: 0.7667 - lr: 0.0010
Epoch 11/50
95/95 [=====] - ETA: 0s - loss: 0.9334 - categorical_accuracy: 0.7123
Epoch 0011: val_loss did not improve from 0.85638
95/95 [=====] - 23s 241ms/step - loss: 0.9334 - categorical_accuracy: 0.7123 - val_loss: 0.8875 - val_categorical_accuracy: 0.6667 - lr: 0.0010
Epoch 12/50
95/95 [=====] - ETA: 0s - loss: 0.9366 - categorical_accuracy: 0.6667
Epoch 0012: val_loss did not improve from 0.85638
95/95 [=====] - 23s 242ms/step - loss: 0.9366 - categorical_accuracy: 0.6667 - val_loss: 0.9344 - val_categorical_accuracy: 0.7000 - lr: 0.0010
Epoch 13/50
95/95 [=====] - ETA: 0s - loss: 0.8739 - categorical_accuracy: 0.7088
Epoch 0013: val_loss did not improve from 0.85638
95/95 [=====] - 24s 255ms/step - loss: 0.8739 - categorical_accuracy: 0.7088 - val_loss: 0.8702 - val_categorical_accuracy: 0.6333 - lr: 0.0010
Epoch 14/50
95/95 [=====] - ETA: 0s - loss: 0.8687 - categorical_accuracy: 0.7298
Epoch 0014: val_loss improved from 0.85638 to 0.58786, saving model to model/MobileNet.h5
95/95 [=====] - 23s 243ms/step - loss: 0.8687 - categorical_accuracy: 0.7298 - val_loss: 0.5879 - val_categorical_accuracy: 0.9333 - lr: 0.0010

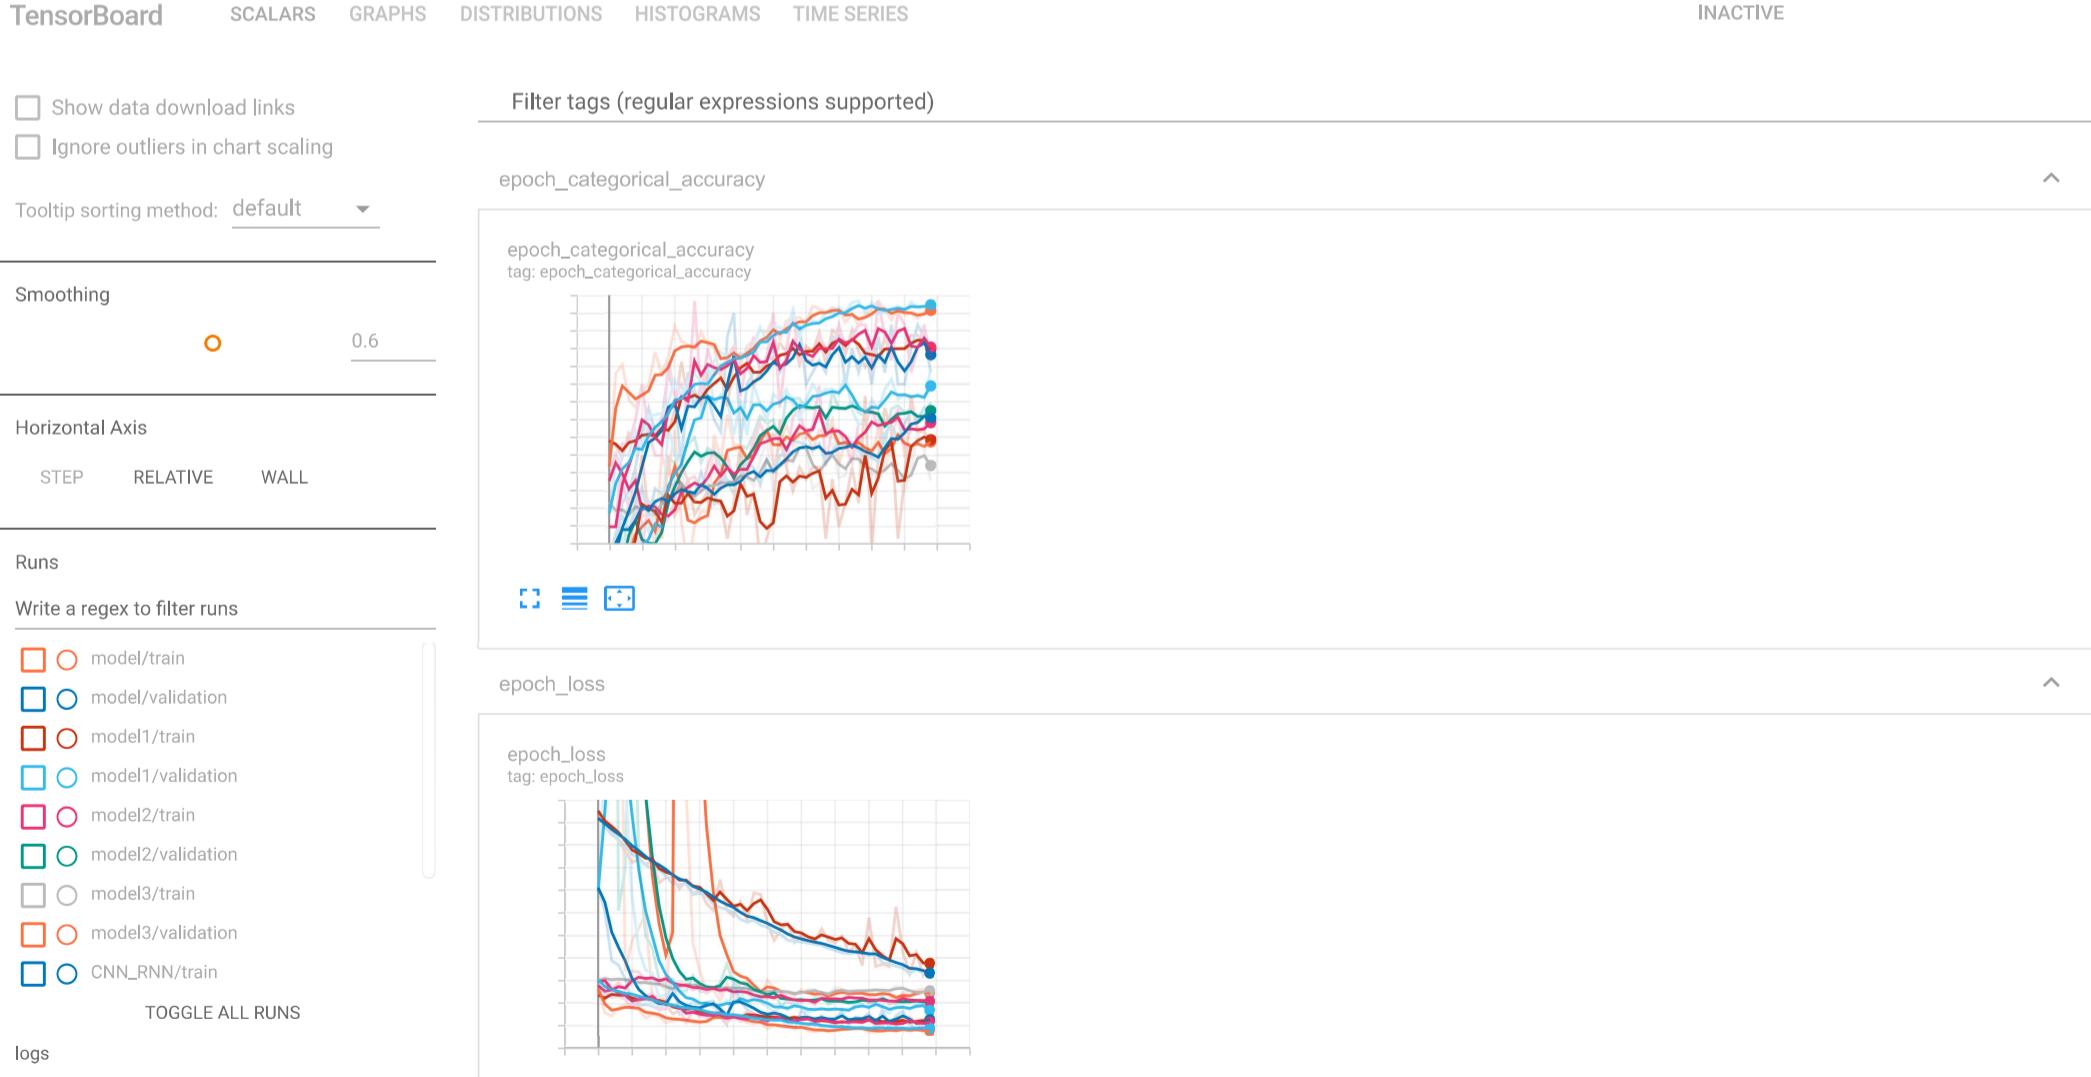
```

1 # Visualizing the Model Results using TensorBoard

2

3 %tensorboard --logdir logs

→ Reusing TensorBoard on port 6006 (pid 646), started 2:05:03 ago. (Use '!kill 646' to kill it.)



Best Epoch Result from the above Model: **Train Accuracy : 93 Validation Accuracy : 90**

We can clearly see that Transfer Learning Model using GRU is performing best.

Transfer Learning Model(Mobile Net : Trainable = False)

```

1 #Changing the batch Size
2 train_generator, val_generator, steps_per_epoch, validation_steps = gen_train_val(batch_size=15)
1 ## Transfer Learning using Mobile Net Architecture
2 mobilenetnt = tf.keras.applications.MobileNetV2(include_top=False)
3 mobilenetnt.trainable = False
4 mobile = Sequential([
5     TimeDistributed(mobilenetnt, input_shape=(30,120,120,3)),
6     TimeDistributed(BatchNormalization()),
7     TimeDistributed(Flatten()),
8     GRU(6,return_sequences=False,dropout=0.25),
9     Dense(5,activation='softmax')
10 ])

```

WARNING:tensorflow:`input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.

```

1 ## Compile the Model
2 optimiser = tf.keras.optimizers.SGD(learning_rate=0.001) # write your optimizer
3 mobile.compile(optimizer=optimiser, loss='categorical_crossentropy', metrics=['CategoricalAccuracy'])
4 mobile.summary()

```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
<hr/>		
time_distributed_17 (TimeDistributed)	(None, 30, 3, 3, 1280)	2257984
time_distributed_18 (TimeDistributed)	(None, 30, 3, 3, 1280)	5120
time_distributed_19 (TimeDistributed)	(None, 30, 11520)	0
gru_2 (GRU)	(None, 6)	207504
dense_14 (Dense)	(None, 5)	35
<hr/>		
Total params: 2,470,643		
Trainable params: 210,099		
Non-trainable params: 2,260,544		

```

1 #Transfer Learning Model fit
2 mobile.fit(train_generator,
3             steps_per_epoch=steps_per_epoch,
4             epochs=50,
5             verbose=1,
6             callbacks=callbacks_list7,
7             validation_data=val_generator,
8             validation_steps=validation_steps)

```

```

Source path = /content/train ; batch size = 15
Epoch 1/50
45/45 [=====] - ETA: 0s - loss: 1.5471 - categorical_accuracy: 0.3363Source path = /content/val ; batch size = 15

Epoch 00001: val_loss improved from inf to 1.56592, saving model to model/MobileNetNonTrainable.h5
45/45 [=====] - 55s 1s/step - loss: 1.5471 - categorical_accuracy: 0.3363 - val_loss: 1.5659 - val_categorical_accuracy: 0.3200 - lr: 0.0010
Epoch 2/50
45/45 [=====] - ETA: 0s - loss: 1.4612 - categorical_accuracy: 0.4222
Epoch 00002: val_loss improved from 1.56592 to 1.32979, saving model to model/MobileNetNonTrainable.h5
45/45 [=====] - 16s 374ms/step - loss: 1.4612 - categorical_accuracy: 0.4222 - val_loss: 1.3298 - val_categorical_accuracy: 0.4429 - lr: 0.0010
Epoch 3/50
45/45 [=====] - ETA: 0s - loss: 1.3197 - categorical_accuracy: 0.4444
Epoch 00003: val_loss improved from 1.32979 to 1.22636, saving model to model/MobileNetNonTrainable.h5
45/45 [=====] - 16s 365ms/step - loss: 1.3197 - categorical_accuracy: 0.4444 - val_loss: 1.2264 - val_categorical_accuracy: 0.4143 - lr: 0.0010
Epoch 4/50
45/45 [=====] - ETA: 0s - loss: 1.2844 - categorical_accuracy: 0.4593
Epoch 00004: val_loss improved from 1.22636 to 1.22389, saving model to model/MobileNetNonTrainable.h5
45/45 [=====] - 15s 350ms/step - loss: 1.2844 - categorical_accuracy: 0.4593 - val_loss: 1.2239 - val_categorical_accuracy: 0.4429 - lr: 0.0010
Epoch 5/50
45/45 [=====] - ETA: 0s - loss: 1.2501 - categorical_accuracy: 0.4741
Epoch 00005: val_loss improved from 1.22389 to 1.18851, saving model to model/MobileNetNonTrainable.h5
45/45 [=====] - 16s 368ms/step - loss: 1.2501 - categorical_accuracy: 0.4741 - val_loss: 1.1885 - val_categorical_accuracy: 0.4571 - lr: 0.0010
Epoch 6/50
45/45 [=====] - ETA: 0s - loss: 1.2013 - categorical_accuracy: 0.5704
Epoch 00006: val_loss improved from 1.18851 to 1.16508, saving model to model/MobileNetNonTrainable.h5
45/45 [=====] - 17s 375ms/step - loss: 1.2013 - categorical_accuracy: 0.5704 - val_loss: 1.1651 - val_categorical_accuracy: 0.4571 - lr: 0.0010
Epoch 7/50
45/45 [=====] - ETA: 0s - loss: 1.0523 - categorical_accuracy: 0.5778
Epoch 00007: val_loss did not improve from 1.16508
45/45 [=====] - 15s 350ms/step - loss: 1.0523 - categorical_accuracy: 0.5778 - val_loss: 1.1831 - val_categorical_accuracy: 0.4429 - lr: 0.0010
Epoch 8/50
45/45 [=====] - ETA: 0s - loss: 1.1176 - categorical_accuracy: 0.6222
Epoch 00008: val_loss improved from 1.16508 to 1.15055, saving model to model/MobileNetNonTrainable.h5

```